

Chapter 3: Feature Extraction

Computer Vision – Unit 03

Computer Vision Course

February 2026

Table of Contents

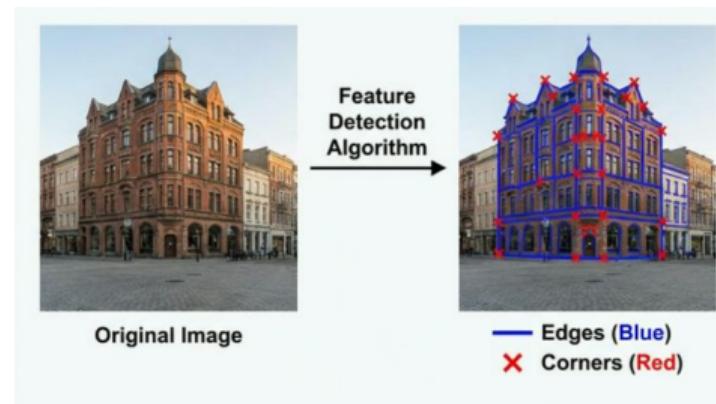
1. Introduction to Feature Extraction
2. Edge Detection
3. Line Detection – Hough Transform
4. Corner Detection
5. Orientation Histogram
6. SIFT – Scale-Invariant Feature Transform
7. SURF – Speeded-Up Robust Features
8. HOG – Histogram of Oriented Gradients
9. GLOH – Gradient Location and Orientation Histogram
10. Scale-Space Analysis
11. Gabor Filters
12. Discrete Wavelet Transform (DWT)
13. Derivation: Harris Corner Detector
14. Numerical Problems
15. Long Numerical Problems
16. Summary

What is Feature Extraction?

- ▶ **Feature extraction** is the process of identifying and isolating meaningful patterns, structures, or descriptors from raw image data.
- ▶ Goal: Reduce data dimensionality while retaining discriminative information.
- ▶ Common features: **Edges, Corners, Blobs, Ridges, Textures.**

Why Feature Extraction?

- ▶ Enables object recognition, tracking, image stitching, 3D reconstruction.
- ▶ Forms the foundation of many computer vision pipelines.



Taxonomy of Features

- ▶ **Edges:** Boundaries between regions (Canny, LOG, DOG)
- ▶ **Lines:** Straight structures (Hough Transform)
- ▶ **Corners / Interest Points:** Harris, Hessian Affine
- ▶ **Descriptors:** SIFT, SURF, HOG, GLOH
- ▶ **Multi-scale:** Scale-Space, Image Pyramids
- ▶ **Frequency-based:** Gabor Filters, DWT

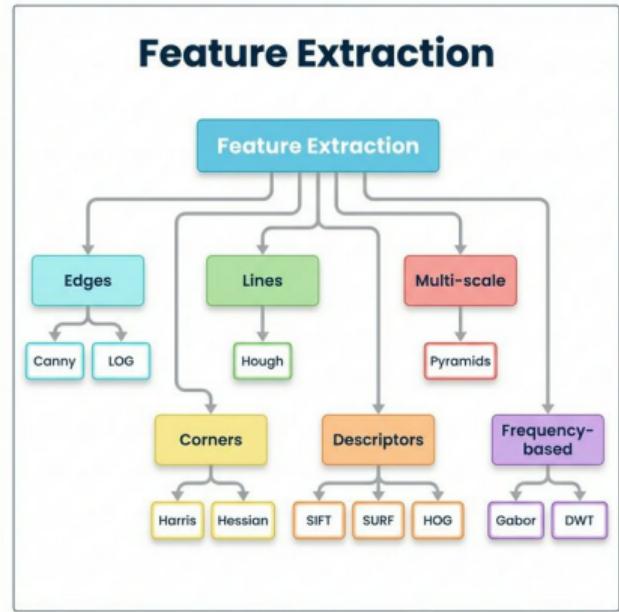


Figure: Hierarchy chart showing feature types branching from “Feature Extraction”

Edge Detection – Overview

- ▶ An **edge** is a significant local change in image intensity.
- ▶ Edges correspond to object boundaries, surface markings, shadows, or depth discontinuities.
- ▶ Mathematically:

$$\nabla I = \left(\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right)$$

- ▶ Magnitude: $|\nabla I| = \sqrt{I_x^2 + I_y^2}$
- ▶ Direction: $\theta = \arctan(I_y/I_x)$

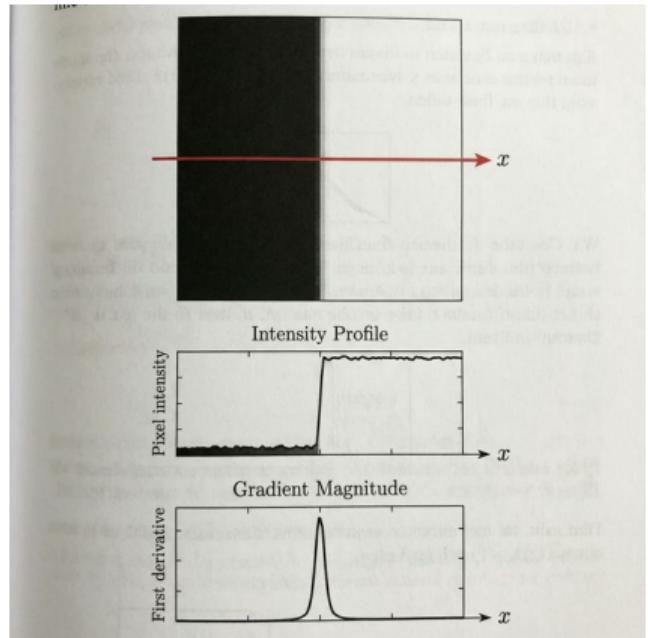


Figure: Image showing intensity profile across an edge with gradient plot

Canny Edge Detector

Canny (1986) – Optimal Edge Detector

Designed to satisfy three criteria: **Good detection**, **Good localisation**, **Single response**.

Steps of the Canny Algorithm:

1. **Gaussian Smoothing:** Convolve with G_σ to reduce noise.

$$G_\sigma(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

2. **Gradient Computation:** Compute $|\nabla I|$ and θ using Sobel or similar operators.
3. **Non-Maximum Suppression (NMS):** Thin edges by keeping only local maxima along gradient direction.
4. **Double Thresholding:** Use T_{high} and T_{low} to classify strong, weak, and non-edges.
5. **Edge Tracking by Hysteresis:** Weak edges connected to strong edges are kept.

Canny Edge Detector – Illustration

Key Parameters:

- ▶ σ – controls smoothing level.
- ▶ $T_{\text{high}}, T_{\text{low}}$ – thresholds for hysteresis.
- ▶ Typical ratio: $T_{\text{high}}/T_{\text{low}} \approx 2 : 1$ or $3 : 1$.

Important

Larger $\sigma \Rightarrow$ fewer edges detected (more smoothing).

Smaller $\sigma \Rightarrow$ more edges but also more noise.

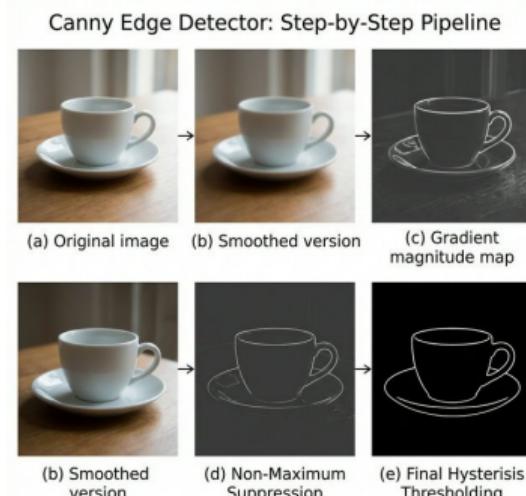


Figure: Step-by-step Canny pipeline – (a) Original (b) Smoothed (c) Gradient (d) NMS (e) Final edges

Laplacian of Gaussian (LOG)

- The **Laplacian** $\nabla^2 I$ is a second-order derivative operator:

$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$

- Direct Laplacian is very sensitive to noise \Rightarrow first smooth with Gaussian.
- **LOG** = Laplacian applied to Gaussian-smoothed image:

$$\text{LOG}(x, y) = \nabla^2 [G_\sigma * I] = [\nabla^2 G_\sigma] * I$$

- The LOG kernel (“Mexican Hat”):

$$\nabla^2 G_\sigma(x, y) = \frac{1}{\pi \sigma^4} \left(\frac{x^2 + y^2}{2\sigma^2} - 1 \right) \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

- Edges are located at **zero-crossings** of the LOG response.

LOG – Visualization

Properties of LOG:

- ▶ Isotropic (rotationally symmetric).
- ▶ Detects edges at zero-crossings.
- ▶ Scale controlled by σ .
- ▶ Sensitive to blobs at scale σ .

Typical 5×5 LOG Mask ($\sigma = 1.4$)

$$\begin{bmatrix} 0 & 0 & -1 & 0 & 0 \\ 0 & -1 & -2 & -1 & 0 \\ -1 & -2 & 16 & -2 & -1 \\ 0 & -1 & -2 & -1 & 0 \\ 0 & 0 & -1 & 0 & 0 \end{bmatrix}$$

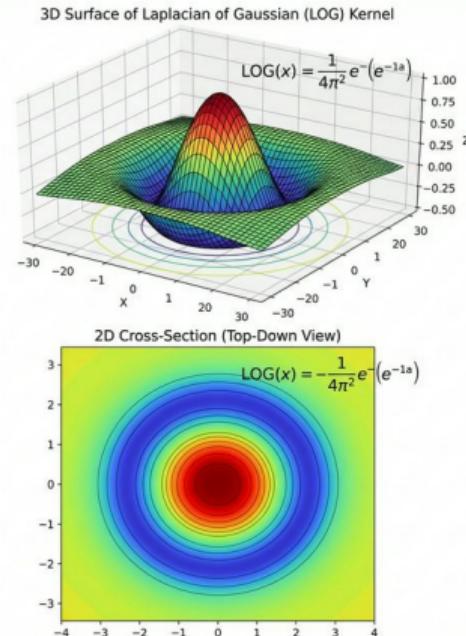


Figure: 3D surface plot of LOG kernel (Mexican Hat shape) and its 2D cross-section

Difference of Gaussians (DOG)

- ▶ **DOG** approximates the LOG and is computationally cheaper:

$$\text{DOG}(x, y) = G_{\sigma_1} - G_{\sigma_2}, \quad \sigma_2 > \sigma_1$$

- ▶ Typically $\sigma_2 = k\sigma_1$ with $k \approx 1.6$.
- ▶ $\text{DOG} \approx \text{LOG}$ because:

$$G_{\sigma_1} - G_{\sigma_2} \approx (\sigma_2 - \sigma_1) \nabla^2 G_{\sigma}$$

- ▶ **Advantages:** Faster computation (subtract two blurred images), directly gives scale information.
- ▶ Used extensively in **SIFT** for keypoint detection.

Difference of Gaussians (DOG) Explained

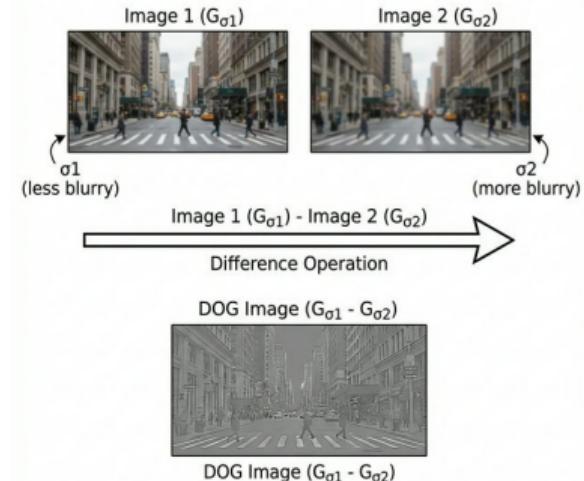


Figure: Two Gaussian-blurred images and their difference showing blob-like edge response

Edge Detectors – Comparison

Property	Canny	LOG	DOG
Derivative Order	1st	2nd	\approx 2nd
Edge Criterion	Gradient max	Zero-crossing	Zero-crossing
Noise Handling	Gaussian smooth	Gaussian smooth	Inherent
Computational Cost	Moderate	High	Low
Thin Edges	Yes (NMS)	Yes	Approximate
Multi-scale	Manual σ	σ param	Natural

Key Takeaway

Canny is the most widely used general-purpose edge detector. LOG gives precise zero-crossings. DOG is a fast approximation used in scale-space frameworks like SIFT.

Hough Transform – Concept

- ▶ **Hough Transform** detects parametric shapes (lines, circles, ellipses) via a **voting scheme** in parameter space.
- ▶ For a line $y = mx + c$, each edge point votes for all (m, c) pairs passing through it.
- ▶ Problem: $m \rightarrow \infty$ for vertical lines.

Normal (Polar) Parameterization

$$x \cos \theta + y \sin \theta = \rho$$

where ρ = perpendicular distance from origin, θ = angle of normal.

Parameter space: (ρ, θ) with $\rho \in [-D, D]$, $\theta \in [0, \pi]$.

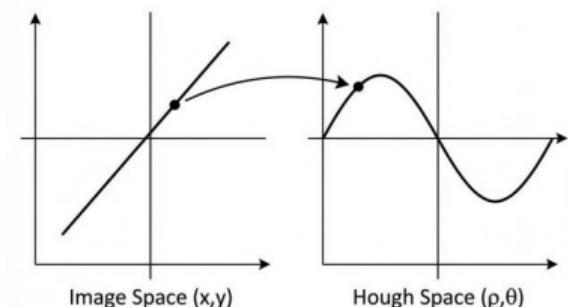


Figure: Line in image space mapped to sinusoidal curve in (ρ, θ) Hough space

Hough Transform – Algorithm

Algorithm Steps:

1. Run edge detection (e.g., Canny) on input image.
2. Create an **accumulator array** $A[\rho][\theta]$ initialized to 0.
3. For each edge pixel (x_i, y_i) :
 - ▶ For each θ from 0 to π (quantized):
 - ▶ Compute $\rho = x_i \cos \theta + y_i \sin \theta$
 - ▶ Increment $A[\rho][\theta]$
4. Find **peaks** in the accumulator \Rightarrow detected lines.

Complexity

$O(n \cdot q_\theta)$ where n = number of edge pixels, q_θ = quantization levels.

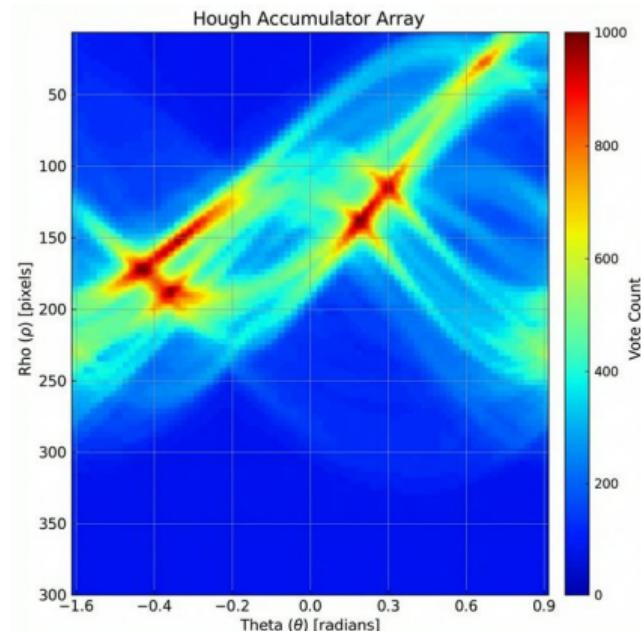


Figure: Accumulator array heatmap with bright peaks at detected line parameters

Hough Transform – Example

Example: Three collinear points $(1, 1)$, $(2, 2)$, $(3, 3)$.

For each point, the Hough curve is:

$$\rho = x \cos \theta + y \sin \theta$$

- ▶ Point $(1, 1)$: $\rho = \cos \theta + \sin \theta$
- ▶ Point $(2, 2)$: $\rho = 2 \cos \theta + 2 \sin \theta$
- ▶ Point $(3, 3)$: $\rho = 3 \cos \theta + 3 \sin \theta$

All three curves intersect at $\theta = 45^\circ$, confirming the line $y = x$.

At $\theta = 45^\circ$: $\rho = \sqrt{2}, 2\sqrt{2}, 3\sqrt{2}$ – Wait, they intersect where all produce same (ρ, θ) , which is at $\theta = 135^\circ, \rho = 0$ (line through origin).

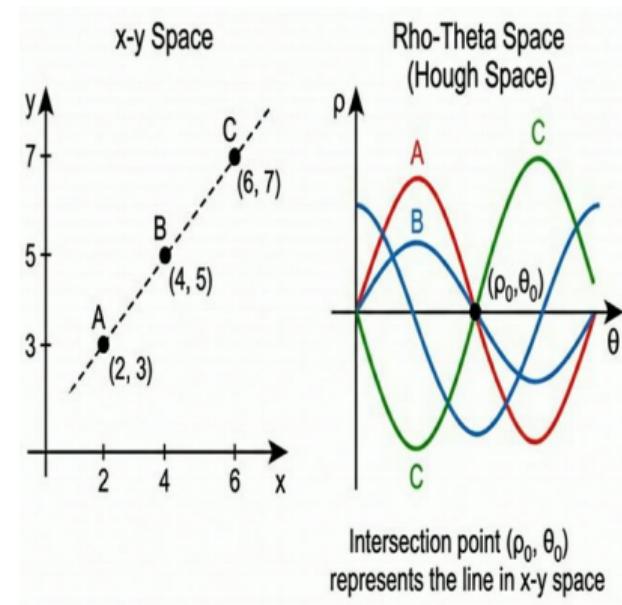


Figure: Sinusoidal Hough curves for three collinear points intersecting at a single (ρ, θ) peak

Harris Corner Detector – Motivation

- ▶ **Corners** are points where intensity changes significantly in *all* directions.
- ▶ Consider a small window shifted by (u, v) . The **Sum of Squared Differences (SSD)**:

$$E(u, v) = \sum_{(x,y) \in W} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

- ▶ Using Taylor expansion $I(x + u, y + v) \approx I(x, y) + I_x u + I_y v$:

$$E(u, v) \approx [u \quad v] \mathbf{M} \begin{bmatrix} u \\ v \end{bmatrix}$$

Structure Tensor (Second Moment Matrix)

$$\mathbf{M} = \sum_{(x,y) \in W} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

Harris Corner Detector – Response Function

- ▶ Let λ_1, λ_2 be eigenvalues of \mathbf{M} :
 - ▶ Flat: $\lambda_1 \approx \lambda_2 \approx 0$
 - ▶ Edge: One $\lambda \gg 0$, other ≈ 0
 - ▶ Corner: Both λ_1, λ_2 large
- ▶ Harris response (no eigenvalues needed):

$$R = \det(\mathbf{M}) - k(\text{trace}(\mathbf{M}))^2$$

where $\det(\mathbf{M}) = \lambda_1 \lambda_2$, $\text{trace}(\mathbf{M}) = \lambda_1 + \lambda_2$,
 $k \in [0.04, 0.06]$.

- ▶ Decision:
 - ▶ $R > T \Rightarrow \text{Corner}$
 - ▶ $R < 0, |R| \text{ large} \Rightarrow \text{Edge}$
 - ▶ $|R| \text{ small} \Rightarrow \text{Flat}$

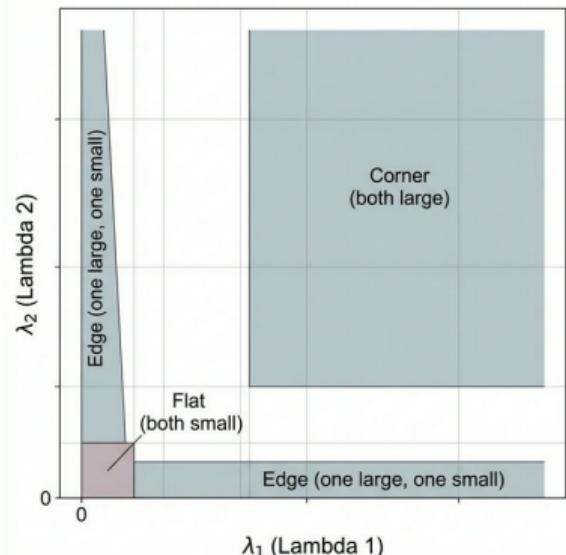


Figure: λ_1 vs λ_2 plot showing corner, edge, and flat regions

Harris Corner Detector – Properties

Algorithm Steps:

1. Compute image gradients I_x, I_y .
2. Compute products $I_x^2, I_y^2, I_x I_y$.
3. Apply Gaussian window to each product.
4. Compute Harris response R at each pixel.
5. Apply non-maximum suppression.
6. Threshold $R > T$ to select corners.

Properties:

- ▶ **Rotation invariant** (eigenvalues unchanged).
- ▶ **NOT scale invariant.**
- ▶ Partially invariant to affine intensity changes.

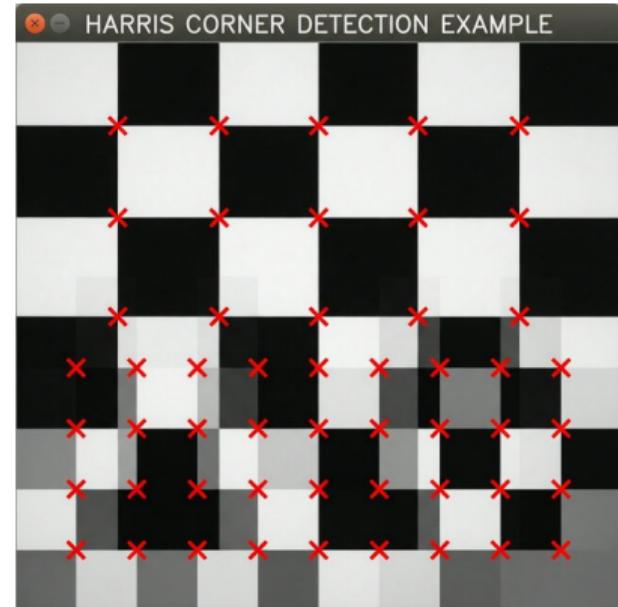


Figure: Image with detected Harris corners marked as red dots overlaid on a checkerboard pattern

Hessian Affine Detector

- The **Hessian matrix** at point (x, y) at scale σ :

$$\mathbf{H}(x, y, \sigma) = \begin{bmatrix} L_{xx}(x, y, \sigma) & L_{xy}(x, y, \sigma) \\ L_{xy}(x, y, \sigma) & L_{yy}(x, y, \sigma) \end{bmatrix}$$

where $L_{xx} = \sigma^2 \frac{\partial^2}{\partial x^2} (G_\sigma * I)$, etc.

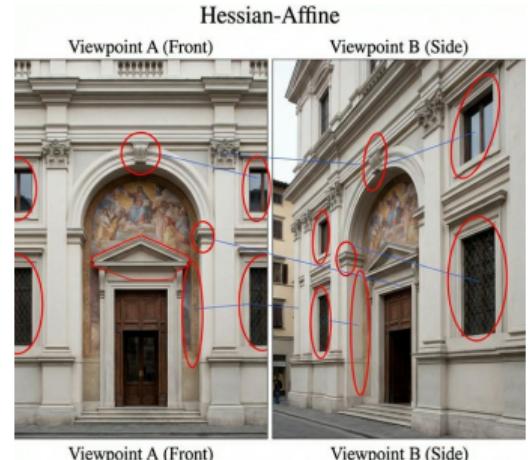
- **Blob/corner detection:** $\det(\mathbf{H}) = L_{xx}L_{yy} - L_{xy}^2 > T$
- **Hessian-Affine** extends this with:
 1. Multi-scale detection using Hessian determinant.
 2. **Affine adaptation:** Iteratively estimate local affine shape using the second moment matrix.
 3. Normalize the region to a canonical circular shape.

Key Advantage

Provides **affine-invariant** interest regions, enabling matching under viewpoint changes.

Harris vs Hessian-Affine – Comparison

Property	Harris	Hessian Affine
Detects	Corners	Blobs + Corners
Matrix	Structure Tensor \mathbf{M}	Hessian \mathbf{H}
Scale Inv.	No	Yes (multi-scale)
Affine Inv.	No	Yes (iterative)
Comp.	Fast	Moderate–High
Repeat.	Good	Very Good



Affine-adapted regions (Hessian-Affine blobs) tracking corresponding surface patches across view changes. Connected lines indicate matches.

Figure: Hessian-Affine regions matched across views

Orientation Histogram

- ▶ An **Orientation Histogram** captures distribution of gradient directions.
- ▶ Steps:
 1. Compute gradient magnitude m and orientation θ .
 2. Quantize θ into n bins (e.g., $n = 36$ bins).
 3. Vote weighted by m .
- ▶ **Dominant orientation** = bin with max count.
- ▶ Used in SIFT \Rightarrow rotation invariance.

Example

If gradients mostly point at 45, the histogram peak is in the [40, 50] bin.

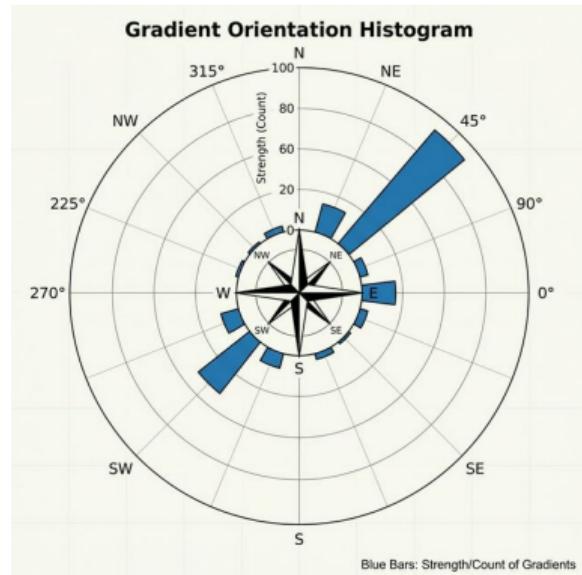


Figure: Polar histogram of gradient orientations with a dominant peak at 45

SIFT – Overview (Lowe, 2004)

Goal

Detect and describe features invariant to **scale**, **rotation**, robust to affine/illumination.

Four Major Steps:

1. **Scale-space extrema**: DOG pyramid.
2. **Keypoint localization**: Sub-pixel.
3. **Orientation assignment**: Histograms.
4. **Keypoint descriptor**: 128-D vector.

SIFT (Scale-Invariant Feature Transform) 4 Major Steps

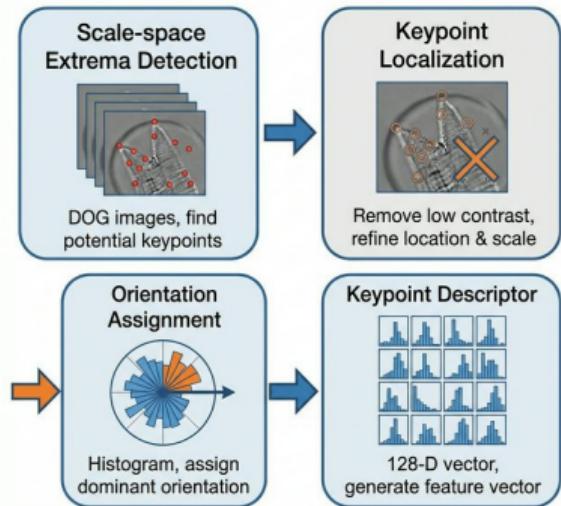


Figure: SIFT pipeline diagram

SIFT – Scale-Space Extrema Detection

- ▶ Build **Gaussian scale space**.
- ▶ Compute **DOG** images.
- ▶ Group into **octaves**.
- ▶ Detect **local extrema**: 26 neighbours.

Extrema Detection in DOG Pyramid

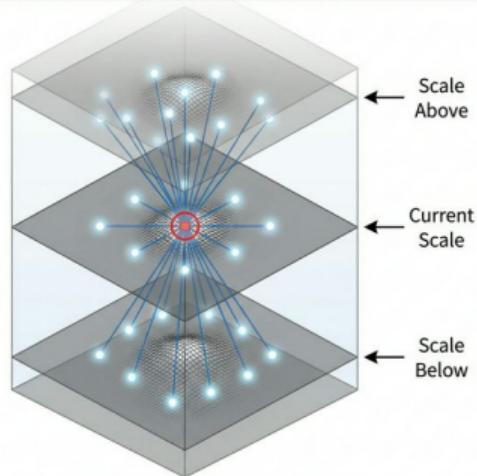


Figure: DOG pyramid extrema

SIFT – Keypoint Localization & Orientation

Keypoint Localization:

- ▶ Fit 3D quadratic (Taylor expansion) to refine position to sub-pixel accuracy:

$$D(\mathbf{x}) \approx D + \frac{\partial D^T}{\partial \mathbf{x}} \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 D}{\partial \mathbf{x}^2} \mathbf{x}$$

- ▶ Reject low-contrast points: $|D(\hat{\mathbf{x}})| < 0.03$.
- ▶ Reject edges using ratio of principal curvatures (Hessian eigenvalue ratio).

Orientation Assignment:

- ▶ Build a 36-bin orientation histogram around each keypoint (weighted by magnitude and Gaussian window).
- ▶ Dominant peak \Rightarrow keypoint orientation.
- ▶ Any peak $\geq 80\%$ of max creates an *additional* keypoint.

SIFT – Descriptor

- ▶ Take 16×16 window, rotate.
- ▶ Divide into 4×4 sub-regions.
- ▶ Compute 8-bin histograms.
- ▶ **128-dim** descriptor.
- ▶ Normalize and clamp.

Matching

Use **Euclidean distance**. Ratio test < 0.8 .

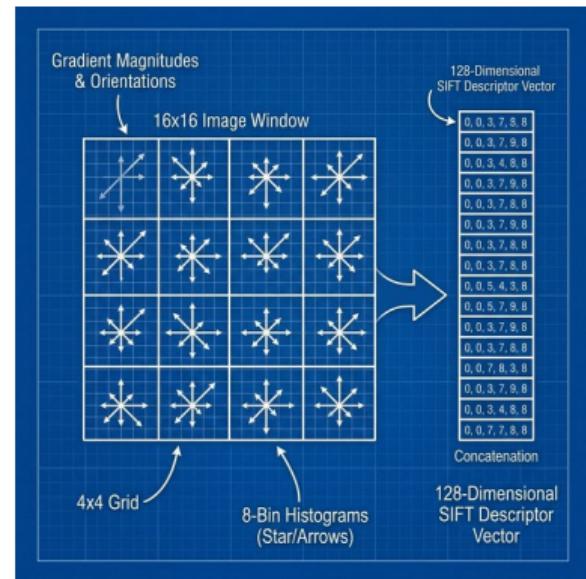


Figure: 128-D SIFT Descriptor

SURF – Overview (Bay et al., 2006)

- ▶ SURF is faster (approximations).
- ▶ Key ideas:
 1. Integral images.
 2. Hessian-based detection.
 3. Haar wavelet responses.

Integral Image

Sum computed in **constant time**.

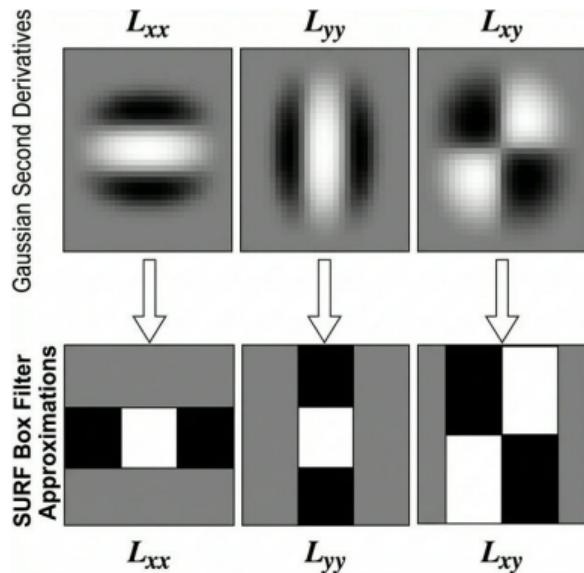


Figure: Box filter approximations

SURF – Detection and Description

Interest Point Detection:

- ▶ Approximate Hessian determinant using box filters at multiple scales.
- ▶ Non-maximum suppression in $3 \times 3 \times 3$ neighbourhood.
- ▶ Sub-pixel interpolation for localization.

Orientation Assignment:

- ▶ Compute Haar wavelet responses in x and y within radius $6s$.
- ▶ Sliding window of $\pi/3$ finds dominant orientation.

Descriptor:

- ▶ 4×4 sub-regions around keypoint.
- ▶ Each sub-region: $(\sum d_x, \sum d_y, \sum |d_x|, \sum |d_y|) \Rightarrow 4$ values.
- ▶ Total: $4 \times 4 \times 4 = \mathbf{64}$ -dimensional descriptor.

SIFT vs SURF – Comparison

Property	SIFT	SURF
Detector	DOG	Hessian (approx)
Dim.	128	64
Speed	Slow	~3× Faster
Scale Inv.	Yes	Yes
Rot. Inv.	Yes	Yes
Robust	Excellent	Good

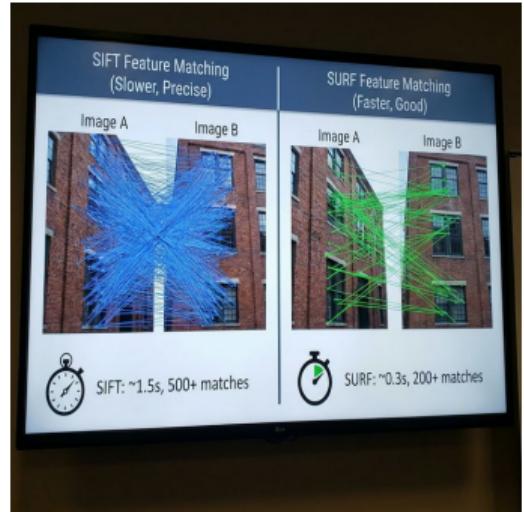


Figure: SIFT (left) vs SURF (right)

HOG – Overview (Dalal & Triggs, 2005)

- ▶ **HOG** describes local shape via distribution of gradient directions.
- ▶ Originally designed for **pedestrian detection**.
- ▶ Key idea: Object appearance can be characterized by the distribution of local intensity gradients, even without precise knowledge of positions.

HOG Pipeline:

1. Compute gradients ($[-1, 0, 1]$ filter).
2. Divide image into **cells** (e.g., 8×8 pixels).
3. Compute 9-bin orientation histogram per cell (unsigned: 0–180).
4. Group cells into overlapping **blocks** (e.g., 2×2 cells).
5. Normalize each block (L2-norm) \Rightarrow illumination invariance.
6. Concatenate all block descriptors.

HOG – Descriptor Construction

Example Calculation:

- ▶ Image window: 64×128 pixels.
- ▶ Cell size: $8 \times 8 \Rightarrow 8 \times 16$ cells.
- ▶ Block size: 2×2 cells with stride 1 cell.
- ▶ Number of blocks: $7 \times 15 = 105$ blocks.
- ▶ Each block: $2 \times 2 \times 9 = 36$ values.
- ▶ Total descriptor: $105 \times 36 = \mathbf{3780}$ dimensions.

Key Properties

- ▶ Captures edge/gradient structure.
- ▶ Block normalization handles illumination.
- ▶ Dense, overlapping computation.

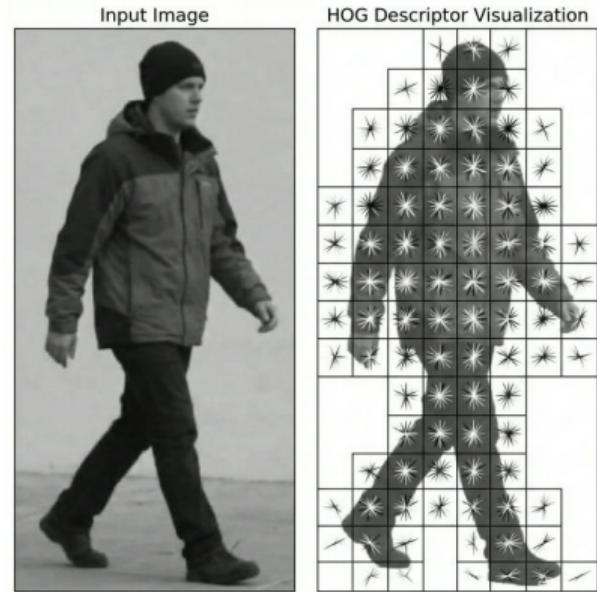


Figure: HOG visualization – image of a person with gradient orientation “hedgehog” glyphs overlaid on cells

GLOH – Gradient Location and Orientation Histogram

- ▶ **GLOH** (Mikolajczyk & Schmid, 2005): SIFT extension.
- ▶ Uses **log-polar grid**:
 - ▶ 3 radial bins (6, 11, 15 px).
 - ▶ 8 angular bins (outer), 1 central.
 - ▶ Total: 17 spatial bins.
- ▶ 16-bin orientation histogram \Rightarrow 272 dims.
- ▶ **PCA** reduction to **128 dims**.

Advantage over SIFT

Log-polar grid better captures radial structure, higher distinctiveness.

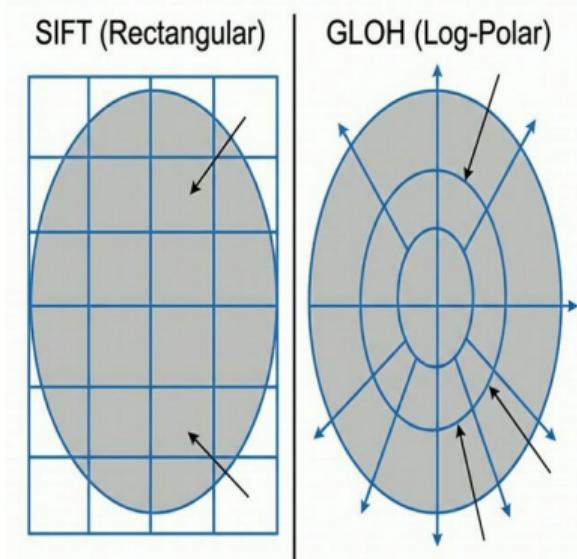


Figure: SIFT (Left) vs GLOH (Right)

Scale-Space Analysis – Motivation

- ▶ Real-world objects exist at **multiple scales**. A feature detector must handle this.
- ▶ **Scale space**: Family of smoothed images $L(x, y, \sigma) = G_\sigma * I(x, y)$ parametrized by scale σ .
- ▶ The Gaussian is the **unique** kernel that satisfies the scale-space axioms (linearity, shift invariance, semi-group, non-enhancement of local extrema).
- ▶ Scale-space representation allows detecting structures at their **characteristic scale**.

Scale-Space Axiom

$$L(x, y, \sigma) = G(x, y; \sigma) * I(x, y), \quad \frac{\partial L}{\partial \sigma} = \sigma \nabla^2 L$$

This is the **diffusion equation** – the scale parameter acts like time in heat diffusion.

Image Pyramids

- ▶ **Image pyramid:** multi-resolution.
- ▶ **Gaussian Pyramid:** Smooth/subsample ($\downarrow 2$).

- ▶ $G_0 = I$
- ▶ $G_{l+1} = \text{subsample}(G_\sigma * G_l)$

- ▶ **Laplacian Pyramid:** Band-pass details.

$$Lap_l = G_l - \text{upsample}(G_{l+1})$$

- ▶ **Overcomplete**, perfect recon.

Applications

Compression, blending, multi-scale.

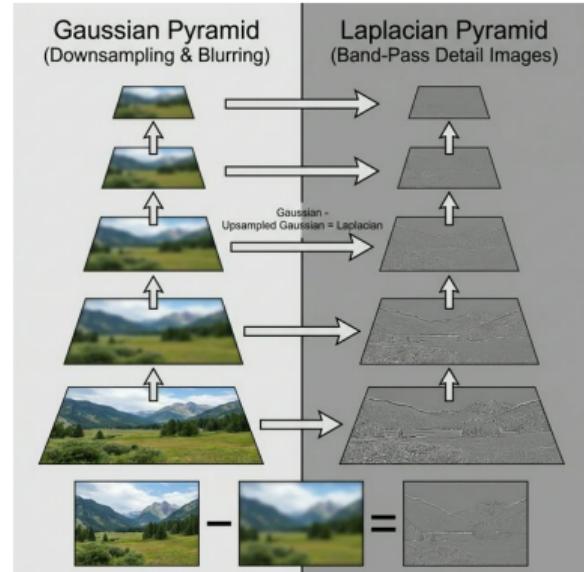


Figure: Gaussian (left) Laplacian (right)

Gaussian Derivative Filters

- ▶ **Gaussian derivative filters** combine smoothing and differentiation in one step:

$$\frac{\partial}{\partial x}(G_\sigma * I) = \left(\frac{\partial G_\sigma}{\partial x} \right) * I$$

- ▶ First-order derivatives of Gaussian:

$$G_x = \frac{\partial G}{\partial x} = -\frac{x}{2\pi\sigma^4} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

- ▶ Second-order: G_{xx} , G_{yy} , G_{xy} – used in Hessian-based detection.
- ▶ These form a **steerable filter** basis: any rotated derivative can be expressed as a linear combination of basis filters.

Scale-Space Feature Detection Summary

- ▶ **Scale selection:** Detect features at the scale where a normalized derivative response is maximal.
- ▶ **Normalized derivatives:** $\sigma^n \nabla^n G_\sigma$ ensures comparable response across scales.
- ▶ **Characteristic scale:** The scale σ^* at which the normalized response achieves its maximum.

$$\sigma^* = \arg \max_{\sigma} |\sigma^2 \nabla^2 L(x, y, \sigma)|$$

Multi-Scale Detection Pipeline

1. Build scale-space (Gaussian pyramid or dense scales).
2. Compute normalized feature response at each scale.
3. Detect extrema across both space *and* scale.
4. Extract features at their characteristic scale.

Gabor Filters – Definition

- ▶ A **Gabor filter** is a Gaussian modulated by a sinusoidal wave:

$$g(x, y; \lambda, \theta, \psi, \sigma, \gamma) = \exp\left(-\frac{x'^2 + \gamma^2 y'^2}{2\sigma^2}\right) \cos\left(2\pi \frac{x'}{\lambda} + \psi\right)$$

where $x' = x\cos\theta + y\sin\theta$, $y' = -x\sin\theta + y\cos\theta$.

- ▶ Parameters:
 - ▶ λ – wavelength of sinusoidal component.
 - ▶ θ – orientation of the filter.
 - ▶ ψ – phase offset.
 - ▶ σ – Gaussian envelope width.
 - ▶ γ – spatial aspect ratio.

Gabor Filters – Properties and Applications

Properties:

- ▶ Achieve optimal **joint resolution** in space and frequency (uncertainty principle).
- ▶ Mimic simple cells in the **human visual cortex**.
- ▶ Selective to **orientation** and **spatial frequency**.

Gabor Filter Bank:

- ▶ Apply filters at multiple orientations (θ) and scales (λ).
- ▶ Typically: 4–8 orientations \times 3–5 scales.
- ▶ Feature vector = concatenated filter responses.

Applications:

- ▶ Texture analysis and segmentation.
- ▶ Face recognition, fingerprint enhancement.
- ▶ Character recognition (OCR).

DWT – Discrete Wavelet Transform

- ▶ DWT decomposes an image into frequency sub-bands at multiple scales.
- ▶ Uses a pair of filters: **low-pass** (scaling function ϕ) and **high-pass** (wavelet ψ).
- ▶ For 2D images, apply row-wise then column-wise:

LL (Approximation)	LH (Horizontal detail)
HL (Vertical detail)	HH (Diagonal detail)

- ▶ **LL** = low-pass in both directions (coarse approximation).
- ▶ **LH, HL, HH** = detail coefficients capturing edges at different orientations.
- ▶ Multi-level DWT: recursively decompose the LL sub-band.

DWT – Haar Wavelet Example

- ▶ Haar wavelet is the simplest wavelet:

$$\psi(x) = \begin{cases} 1 & 0 \leq x < 0.5 \\ -1 & 0.5 \leq x < 1 \\ 0 & \text{otherwise} \end{cases}$$

- ▶ Low-pass filter: $h = \frac{1}{\sqrt{2}}[1, 1]$, High-pass filter: $g = \frac{1}{\sqrt{2}}[1, -1]$.

1D Haar DWT Example

Signal: [4, 6, 10, 12]

Averages (LL): $\frac{4+6}{2} = 5$, $\frac{10+12}{2} = 11 \Rightarrow [5, 11]$

Differences (HL): $\frac{4-6}{2} = -1$, $\frac{10-12}{2} = -1 \Rightarrow [-1, -1]$

DWT coefficients: [5, 11 | -1, -1]

Applications in CV: Texture classification, image compression (JPEG 2000), feature extraction for recognition.

Gabor vs DWT – Comparison

Property	Gabor Filters	DWT
Basis	Non-orthogonal	Orthogonal
Redundancy	High (overcomplete)	Minimal (critical sampling)
Orientation Selectivity	Continuous θ	Fixed (H, V, D)
Frequency Resolution	Tunable λ	Dyadic scales
Reconstruction	Approximate	Perfect
Computation	Moderate	Fast (pyramidal)
Bio-inspired	Yes (V1 cortex)	No

When to Use

Gabor: When fine orientation and frequency tuning is needed (face, texture).

DWT: When compact, non-redundant multi-scale representation is needed (compression, fast features).

Derivation – Harris Corner Response Function

Goal: Derive the Harris corner response R from the auto-correlation function.

Step 1: Define the weighted SSD for a shift (u, v) :

$$E(u, v) = \sum_{(x,y)} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Step 2: Apply first-order Taylor expansion:

$$I(x + u, y + v) \approx I(x, y) + I_x u + I_y v$$

$$\Rightarrow E(u, v) \approx \sum_{(x,y)} w(x, y) [I_x u + I_y v]^2$$

Step 3: Expand the square:

$$E(u, v) \approx \sum w(x, y) [I_x^2 u^2 + 2I_x I_y uv + I_y^2 v^2]$$

Derivation – Harris Corner Response (cont.)

Step 4: Write in matrix form:

$$E(u, v) \approx [u \quad v] \underbrace{\sum_{(x,y)} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}}_{\mathbf{M}} \begin{bmatrix} u \\ v \end{bmatrix}$$

Step 5: Analyze eigenvalues λ_1, λ_2 of \mathbf{M} :

- ▶ $E(u, v)$ is an ellipse with axes λ_1, λ_2 .
- ▶ Both large \Rightarrow corner; one large \Rightarrow edge; both small \Rightarrow flat.

Step 6: Define Harris response (avoids eigenvalue computation):

$$R = \det(\mathbf{M}) - k[\text{trace}(\mathbf{M})]^2 = \lambda_1 \lambda_2 - k(\lambda_1 + \lambda_2)^2$$

where $k \in [0.04, 0.06]$. **Corner:** $R > 0$ (large); **Edge:** $R < 0$; **Flat:** $|R| \approx 0$. \square

Numerical 1: Gradient at a Pixel

Problem

Given a 3×3 image patch: $\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 3 \\ 1 & 1 & 1 \end{bmatrix}$.

Find I_x and I_y at center pixel using: $I_x = I(\text{right}) - I(\text{left})$, $I_y = I(\text{bottom}) - I(\text{top})$.

Solution:

- ▶ Center pixel value = 1
- ▶ $I_x = I(\text{right}) - I(\text{left}) = 3 - 1 = \mathbf{2}$
- ▶ $I_y = I(\text{bottom}) - I(\text{top}) = 1 - 1 = \mathbf{0}$
- ▶ Magnitude: $|\nabla I| = \sqrt{2^2 + 0^2} = \sqrt{4} = \mathbf{2}$
- ▶ Direction: $\theta = \arctan(0/2) = \arctan(0) = \mathbf{0}$

Meaning

Gradient is purely horizontal ($\theta = 0$), so there is a **vertical edge** to the right of the center pixel.

Numerical 2: Harris Corner Response

Problem

Given: $\mathbf{M} = \begin{bmatrix} 5 & 0 \\ 0 & 5 \end{bmatrix}$, $k = 0.04$. Find R and classify.

Solution:

- ▶ $\det(\mathbf{M}) = 5 \times 5 - 0 \times 0 = \mathbf{25}$
- ▶ $\text{trace}(\mathbf{M}) = 5 + 5 = \mathbf{10}$
- ▶ $R = \det - k \cdot (\text{trace})^2 = 25 - 0.04 \times 100 = 25 - 4 = \boxed{21}$
- ▶ $R > 0 \Rightarrow \mathbf{Corner} \checkmark$

If instead $\mathbf{M} = \begin{bmatrix} 10 & 0 \\ 0 & 0 \end{bmatrix}$:

- ▶ $\det = 0$, $\text{trace} = 10$, $R = 0 - 0.04 \times 100 = -4 \Rightarrow \mathbf{Edge}$

If instead $\mathbf{M} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix}$:

- ▶ $R = 0 \Rightarrow \mathbf{Flat \ region}$

Numerical 3: Hough Transform – Compute ρ

Problem

An edge point is at $(x, y) = (1, 0)$. Compute ρ for $\theta = 0$ and $\theta = 90$.

Solution:

Formula: $\rho = x \cos \theta + y \sin \theta$

- At $\theta = 0$:

$$\rho = 1 \cdot \cos 0 + 0 \cdot \sin 0 = 1 \times 1 + 0 = \boxed{1}$$

- At $\theta = 90$:

$$\rho = 1 \cdot \cos 90 + 0 \cdot \sin 90 = 1 \times 0 + 0 = \boxed{0}$$

Meaning

At $\theta = 0$, the line is vertical at $x = 1$. At $\theta = 90$, the line is horizontal passing through the origin. Each edge point traces a curve in (ρ, θ) space.

Numerical 4: SIFT & SURF Descriptor Size

Problem

Calculate the SIFT and SURF descriptor sizes.

SIFT:

- ▶ Sub-regions: $4 \times 4 = 16$
- ▶ Bins per sub-region: 8
- ▶ Total = $16 \times 8 = \boxed{128}$ dimensions

SURF:

- ▶ Sub-regions: $4 \times 4 = 16$
- ▶ Values per sub-region: 4 ($\sum d_x, \sum d_y, \sum |d_x|, \sum |d_y|$)
- ▶ Total = $16 \times 4 = \boxed{64}$ dimensions

Key Takeaway

SURF descriptor is **half** the size of SIFT \Rightarrow faster to compute and match, but slightly less distinctive.

Numerical 5: Image Pyramid Sizes

Problem

An image is 128×128 . Build a 4-level Gaussian pyramid (subsample by 2). What is the size at each level?

Solution:

Level	Size	Pixels
0	128×128	16,384
1	64×64	4,096
2	32×32	1,024
3	16×16	256

$$\text{Total pixels} = 16384 + 4096 + 1024 + 256 = \mathbf{21,760}$$

$$\text{Overhead} = \frac{21760}{16384} \approx 1.33 \times \text{original} \Rightarrow \text{only } 33\% \text{ extra memory!}$$

Numerical 6: Haar DWT (4 values)

Problem

Apply 1 level of Haar DWT to $f = [6, 2, 8, 4]$.

Solution:

Averages (approximation):

$$\blacktriangleright a_1 = \frac{6+2}{2} = 4, \quad a_2 = \frac{8+4}{2} = 6$$

Differences (detail):

$$\blacktriangleright d_1 = \frac{6-2}{2} = 2, \quad d_2 = \frac{8-4}{2} = 2$$

Output: $[4, 6 | 2, 2]$

Reconstruct:

$$a_1 + d_1 = 4 + 2 = 6 \checkmark, \quad a_1 - d_1 = 4 - 2 = 2 \checkmark$$

$$a_2 + d_2 = 6 + 2 = 8 \checkmark, \quad a_2 - d_2 = 6 - 2 = 4 \checkmark$$

Long Numerical 1: Canny Edge Detection (Step-by-Step)

Problem

Given a 5×5 image. After Gaussian smoothing, the center 3×3 gradient magnitudes and directions are:

Gradient Magnitude:

$$\begin{bmatrix} 2 & 8 & 3 \\ 1 & 10 & 4 \\ 2 & 7 & 3 \end{bmatrix}$$

Gradient Direction:

$$\begin{bmatrix} 45 & 90 & 45 \\ 0 & 90 & 0 \\ 135 & 90 & 135 \end{bmatrix}$$

Apply: (a) Non-Maximum Suppression (NMS), (b) Double thresholding with $T_H = 7$, $T_L = 3$.

Step 1 – NMS at center pixel (magnitude = 10, direction = 90):

Compare along gradient direction (vertical): neighbours are 8 (top) and 7 (bottom).

$10 > 8$ and $10 > 7 \Rightarrow$ **keep** center pixel (it is a local max).

Long Numerical 1: Canny (cont.)

Step 1 (cont.) – NMS for all pixels:

Pixel	Mag	Dir	After NMS
(0, 0): mag=2	45	compare diag: 10, –	suppressed ($2 < 10$)
(0, 1): mag=8	90	compare vert: –, 10	suppressed ($8 < 10$)
(1, 0): mag=1	0	compare horiz: –, 10	suppressed ($1 < 10$)
(1, 1): mag=10	90	compare: 8, 7	kept ($10 >$ both)
(2, 1): mag=7	90	compare: 10, –	suppressed ($7 < 10$)

Step 2 – Thresholding ($T_H = 7$, $T_L = 3$):

- Center pixel: $\text{mag} = 10 > T_H = 7 \Rightarrow \text{Strong edge } \checkmark$
- All other pixels were suppressed in NMS \Rightarrow not edges.

Result: Only the center pixel is detected as a **strong edge point**.

Long Numerical 2: Harris Corner Detection (Full Example)

Problem

A 3×3 image window has pixel values: $I = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 5 \\ 1 & 5 & 5 \end{bmatrix}$.

Compute Harris response at center pixel with $k = 0.05$ (use simple differences, no Gaussian weighting for simplicity).

Step 1 – Compute gradients at each pixel (using right-left, bottom-top):

I_x (right – left):

Only center has both neighbours:

$$I_x = I(1, 2) - I(1, 0) = 5 - 1 = 4$$

I_y (bottom – top):

$$I_y = I(2, 1) - I(0, 1) = 5 - 1 = 4$$

Step 2 – Form structure tensor \mathbf{M} at center (single pixel, no window sum needed):

$$\mathbf{M} = \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} = \begin{bmatrix} 16 & 16 \\ 16 & 16 \end{bmatrix}$$

Long Numerical 2: Harris Corner (cont.)

Step 3 – Compute $\det(\mathbf{M})$ and $\text{trace}(\mathbf{M})$:

- ▶ $\det(\mathbf{M}) = 16 \times 16 - 16 \times 16 = 256 - 256 = \mathbf{0}$
- ▶ $\text{trace}(\mathbf{M}) = 16 + 16 = \mathbf{32}$

Step 4 – Compute Harris response:

$$R = \det(\mathbf{M}) - k \cdot [\text{trace}(\mathbf{M})]^2 = 0 - 0.05 \times 32^2 = 0 - 0.05 \times 1024$$

$$R = -51.2$$

Step 5 – Classification:

- ▶ $R < 0$ (negative) ⇒ **Edge point**, not a corner.

Interpretation

Since $I_x = I_y = 4$ and both are equal, the gradient points in one direction (45). The eigenvalues are $\lambda_1 = 32, \lambda_2 = 0$, meaning intensity changes along only one direction ⇒ this is an **edge**, not a corner.

Chapter Summary

1. **Edge Detection:** Canny (optimal multi-step), LOG (zero-crossings), DOG (fast approx.).
2. **Line Detection:** Hough Transform – voting in (ρ, θ) parameter space.
3. **Corner Detection:** Harris (structure tensor \mathbf{M} , response R); Hessian Affine (scale + affine invariant).
4. **Orientation Histogram:** Gradient direction distribution \Rightarrow rotation invariance.
5. **SIFT:** 128-D descriptor, scale & rotation invariant (DOG + gradient histograms).
6. **SURF:** 64-D, faster SIFT alternative (integral images + Haar wavelets).
7. **HOG:** Dense gradient histograms with block normalization (pedestrian detection).
8. **GLOH:** Log-polar SIFT variant + PCA for compactness.
9. **Scale-Space:** Gaussian pyramids, derivative filters for multi-scale analysis.
10. **Gabor & DWT:** Frequency-domain features for texture analysis.

References

-  J. Canny, "A Computational Approach to Edge Detection," *IEEE TPAMI*, 1986.
-  C. Harris and M. Stephens, "A Combined Corner and Edge Detector," *Alvey Vision Conf.*, 1988.
-  D. Lowe, "Distinctive Image Features from Scale-Invariant Keypoints," *IJCV*, 2004.
-  H. Bay et al., "SURF: Speeded Up Robust Features," *ECCV*, 2006.
-  N. Dalal and B. Triggs, "Histograms of Oriented Gradients for Human Detection," *CVPR*, 2005.
-  K. Mikolajczyk and C. Schmid, "A Performance Evaluation of Local Descriptors," *IEEE TPAMI*, 2005.
-  R.C. Gonzalez and R.E. Woods, *Digital Image Processing*, Prentice Hall.
-  R. Szeliski, *Computer Vision: Algorithms and Applications*, Springer.